③ The Model Optimizer

- Pre-trained models
- The Model Optimizers
- The Inference Engine
- Deploying the app

Basically, a model optimizer is a converter that converts models into an intermediate representation (IR), that can then be fed to an inference engine.

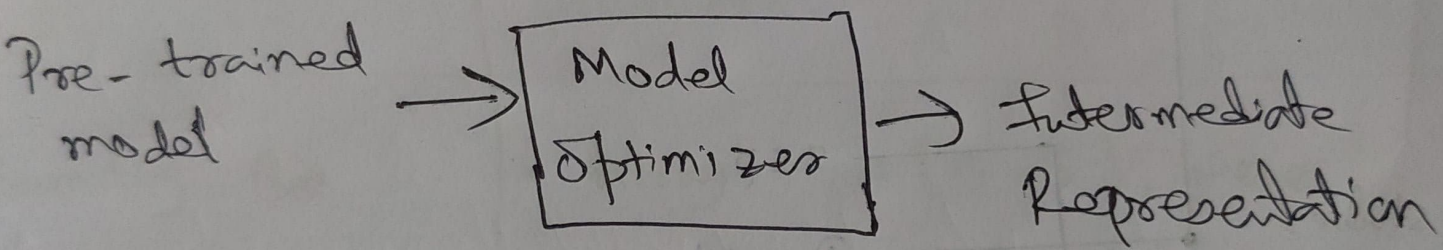

Improvements:

- Model size
- Speed

Trade-offs

- Accuracy (minimized)

Note: Using a pre-trained model is mandatory.

Optimization Techniques

① Quantization

- Refers to the number of bits used to represent the weights and biases of the model.

Model before optimizatn	Model after optimizatn
↑ accuracy	↓ accuracy (not substantial)
↑ size	↓ size
↓ compute ~~time~~ speed	↑ compute ~~time~~ speed

- Model precisions

	Default precision	Present in Open VINO	Supported by Model Optimizer
~~FP32~~		✓	✓
FP32	✓	✓	✓
FP16	X	✓	✓
INT8	X	✓	X

② Freezing

- Used in context of Tensorflow models.
- Removes operations and meta data only needed for training and not inference

Eg: backpropagation.

③ Fusion

- Combines multiple layer operations into a single layer.

Eg:

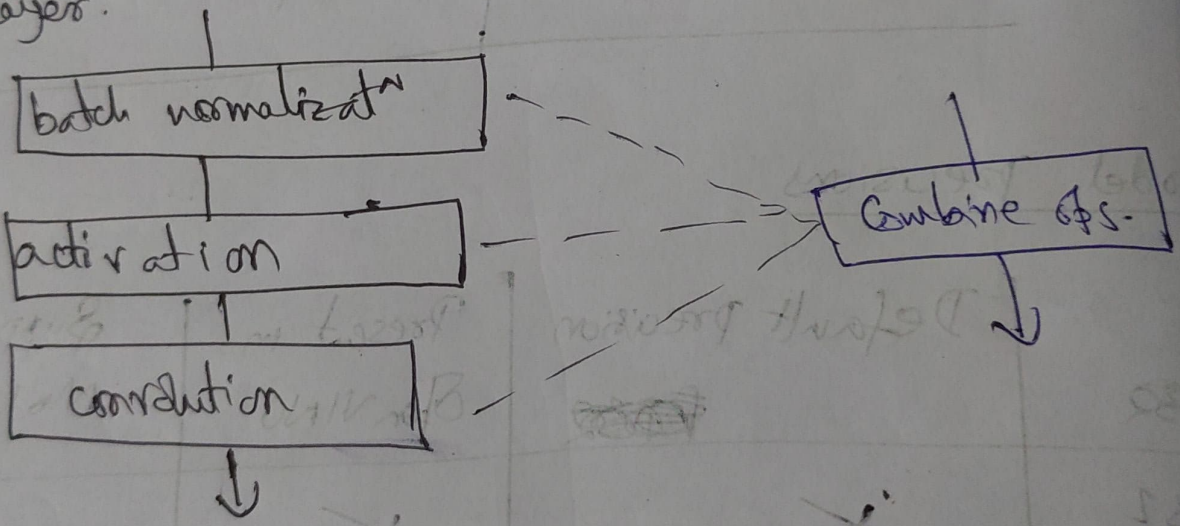

- Useful when diff. operations perform in diff. kernels, but fused operation performs in one kernel

$\Leftrightarrow$ Overhead decreases while switching from one kernel to next

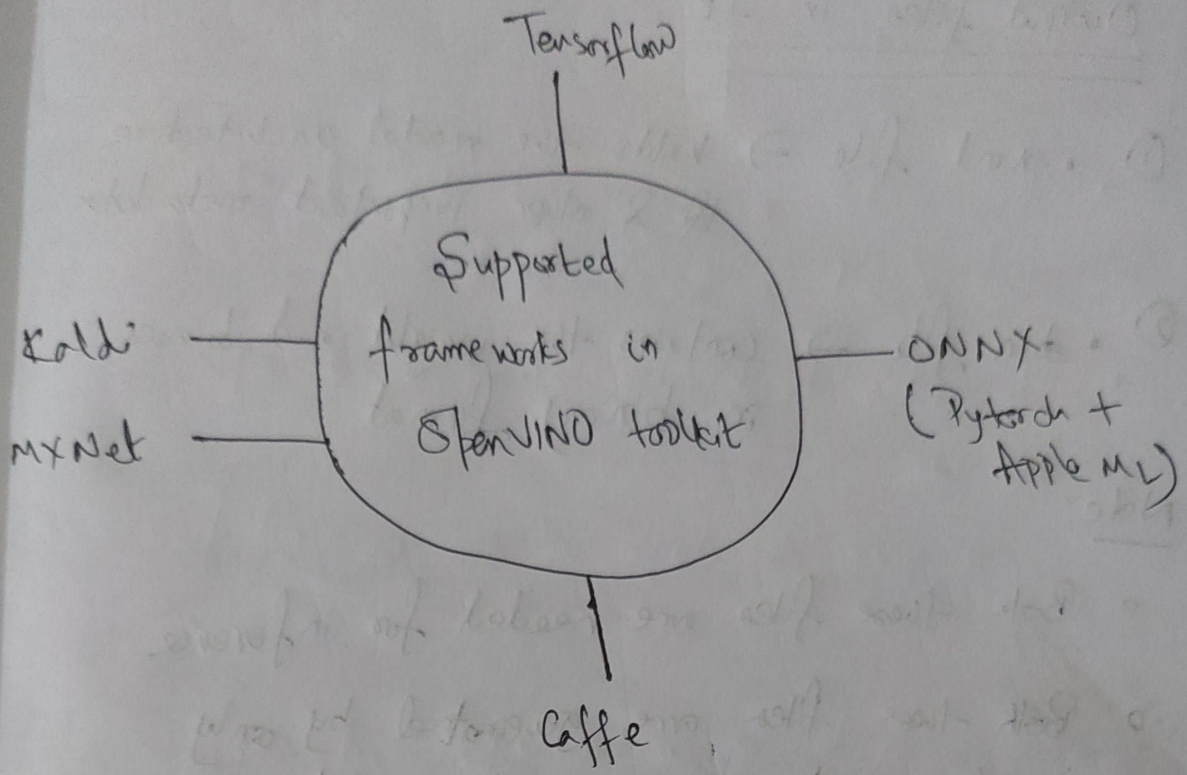

Intermediate Representations

- Refers to the standard structure and naming of neural network architectures in OpenVINO toolkit.

⇔ like "shared dialect" for neural network layers across all the supported frameworks.

Eg:

Conv2D (tensorflow)	Convolution (IR)
Convolution (caffe)	
Conv (ONNX)	

- The model optimizer does this translation into a shared dialect that the Inference Engine can understand.

Output files in IR

1. .xml file → holds the model architecture & other important metadata
2. .bin file → contains the weight and biases in binary format

Note:

- Both these files are needed for inference
- Both these files are generated by only specifying the precision to Model Optimizer using --data_type flag.
  - (By default, the data type was FP32 precision)
- For unfrozen model, use --mean_values and --scale (in case of tensorflow).
- for tensorflow models, remember to:
  - i) reverse i/p channels (BGR → RGB)
  - ii) use custom operations config
  - iii) apply object detection API pipeline conf

Cutting Parts of a Model

Why

- Model has pre- / post-processing parts that don't translate to existing IR layers.
- Model has training part which is not needed during inference.
- Model is too complex w/ many unsupported applications.
- Model is one of the SSD models
  ⇒ cut off the post-processing part
- Cutting the model helps localize any issue in model conversion

How

- CLI args:
  i) --input: if cutting from beginning
  ii) --output: if cutting from end

Custom Layers

- The model may have unsupported layers that need to be handled.
- In order to handle unsupported layers, we can:
  - run the layer w/ the original framework,
    OR
  - write a custom layer that supports the model optimizer

Adding Custom Layers

Tensorflow	Caffe
1. Register the custom layers as extensions to the model optimizer	Register the custom layers as extensions to the model optimizer
2. Replace the unsupported subgraph w/ a different subgraph	Register the layer as a Custom
3. Offload the computation of the subgraph back to TF during inference.	Use Caffe to calculate the o/p of the shape of the layer.

Summary

- **Model Optimizer**
  - Improvements: model size (↓), speed (↑)
  - Tradeoffs: accuracy (↓, minimized)
- **Optimization techniques**
  - Quantization
  - freezing
  - Fusion
- **Supported frameworks in OpenVINO toolkit**
  - Tensorflow
  - Caffe
  - ONNX (Pytorch + Apple ML)
  - MXNet
  - Kaldi
- **Intermediate Representation (IR)**
  - .xml
  - .bin
- **Partial Model Cutting**
  - `--input`
  - `--output`
- **Handling Unsupported Layers**
  - run the layer w/ original framework
  - write a custom layer that supports the model optimizer